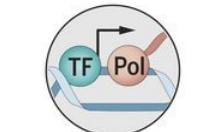


Artificial Neural Networks for MultiOmics Integration

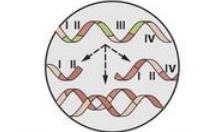
Physalia course, online via zoom

Nikolay Oskolkov, MRG Group Leader, LIOS, Riga, Latvia

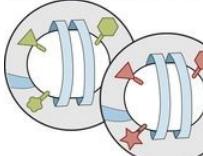


Transcription factor binding

TF binding interacts with DNA methylation and chromatin accessibility



Transcription and RNA maturation



Histone modifications

Modifications can be active marks (e.g., H3K4me3 in green) or repressive marks (e.g., H2K27m3 in red)



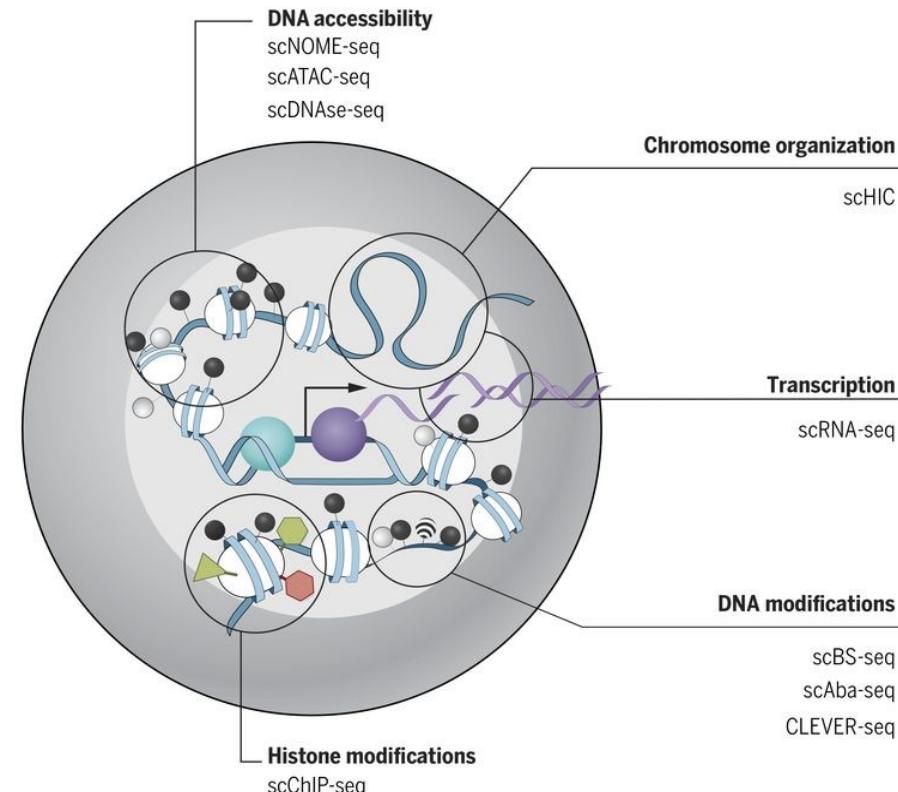
DNA modifications

C ● 5mC
⌚ 5hmC / 5fC / 5caC



Chromosome organization

Higher-order chromatin organization into LADs and TADs



@NikolayOskolkov

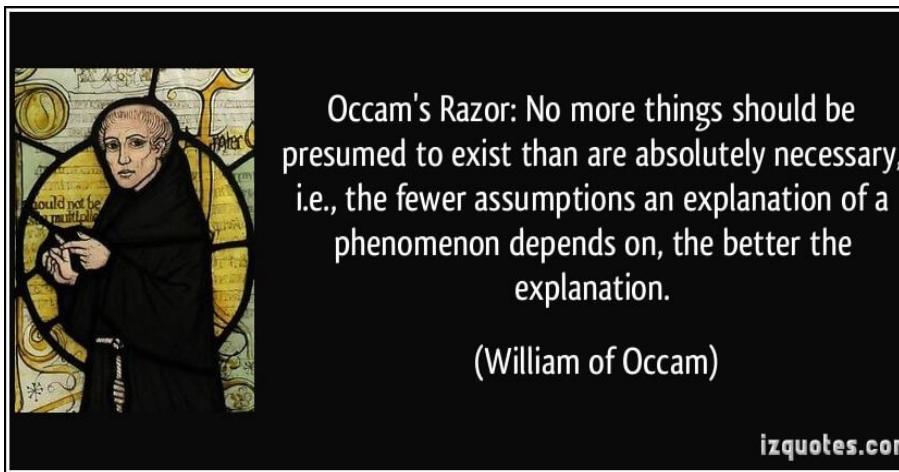


@oskolkov.bsky.social



Personal homepage:
<https://nikolay-oskolkov.com>

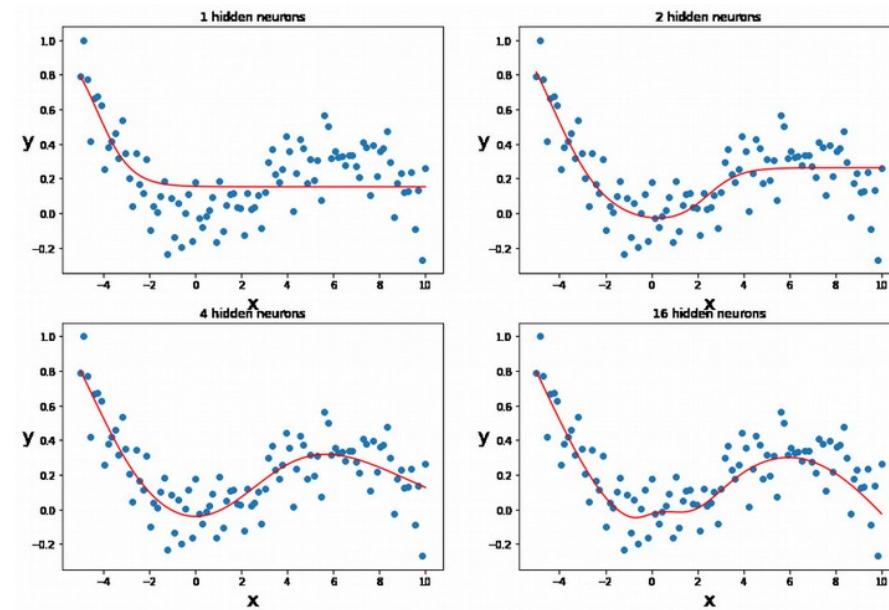
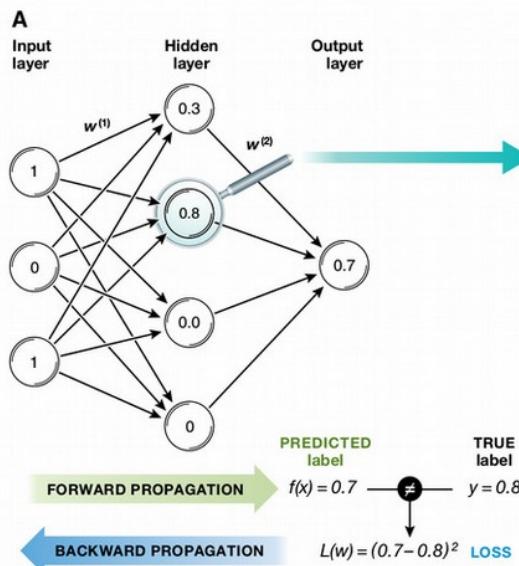
- Difficult to apply to real Life Science projects (NGS: tabular data)
- Lack of data in Life Sciences (exceptions: single cell, microscopy)
- Simpler (than Deep Learning) methods often perform better

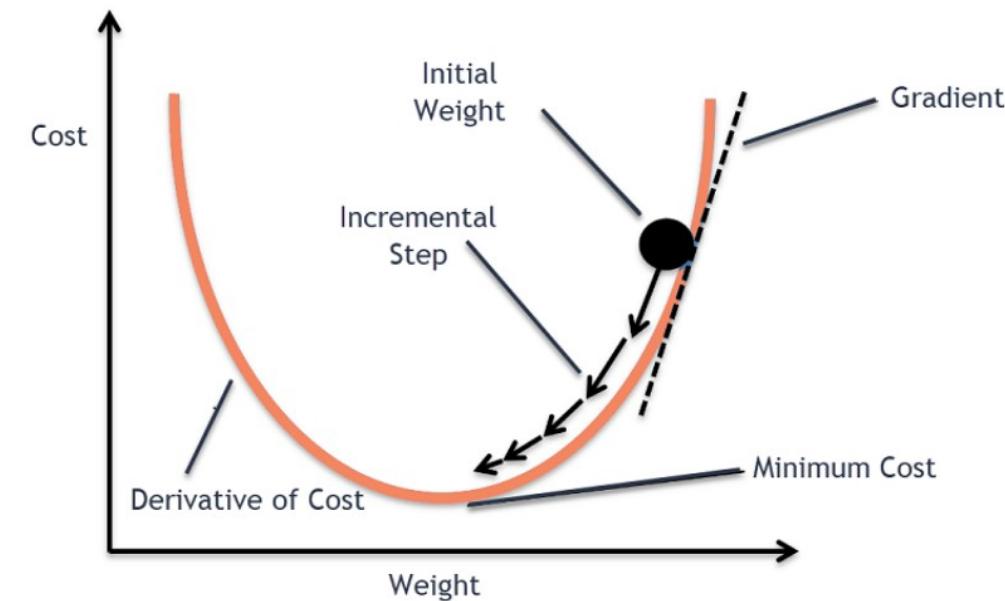


Why don't neural networks always work?

- ANN: a mathematical function $Y = f(X)$ with a special architecture
- Can be non-linear depending on activation function

- Backward propagation (gradient descent) for minimizing error
- Universal Approximation Theorem





$$y_i = \alpha + \beta x_i + \epsilon, \quad i = 1 \dots n$$

$$E(\alpha, \beta) = \frac{1}{n} \sum_{i=1}^n (y_i - \alpha - \beta x_i)^2$$

$$\hat{\alpha}, \hat{\beta} = \operatorname{argmin} E(\alpha, \beta)$$

$$\frac{\partial E(\alpha, \beta)}{\partial \alpha} = -\frac{2}{n} \sum_{i=1}^n (y_i - \alpha - \beta x_i)$$

$$\frac{\partial E(\alpha, \beta)}{\partial \beta} = -\frac{2}{n} \sum_{i=1}^n x_i (y_i - \alpha - \beta x_i)$$

Numeric implementation of gradient descent:

$$\alpha_{i+1} = \alpha_i - \eta \left. \frac{\partial E(\alpha, \beta)}{\partial \alpha} \right|_{\alpha=\alpha_i, \beta=\beta_i}$$

$$\beta_{i+1} = \beta_i - \eta \left. \frac{\partial E(\alpha, \beta)}{\partial \beta} \right|_{\alpha=\alpha_i, \beta=\beta_i}$$

```

1 n <- 100 # sample size
2 x <- rnorm(n) # simulated explanatory variable
3 y <- 3 + 2 * x + rnorm(n) # simulated response variable
4 summary(lm(y ~ x))

```

Call:
`lm(formula = y ~ x)`

Residuals:

Min	1Q	Median	3Q	Max
-1.9073	-0.6835	-0.0875	0.5806	3.2904

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	2.89720	0.09755	29.70	<2e-16 ***
x	1.94753	0.10688	18.22	<2e-16 ***

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

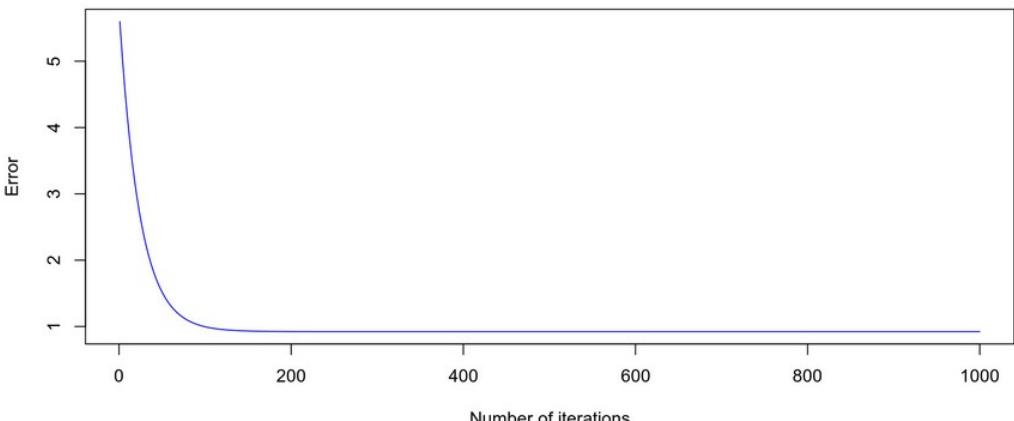
Residual standard error: 0.9707 on 98 degrees of freedom
Multiple R-squared: 0.7721, Adjusted R-squared: 0.7698
F-statistic: 332 on 1 and 98 DF, p-value: < 2.2e-16

```

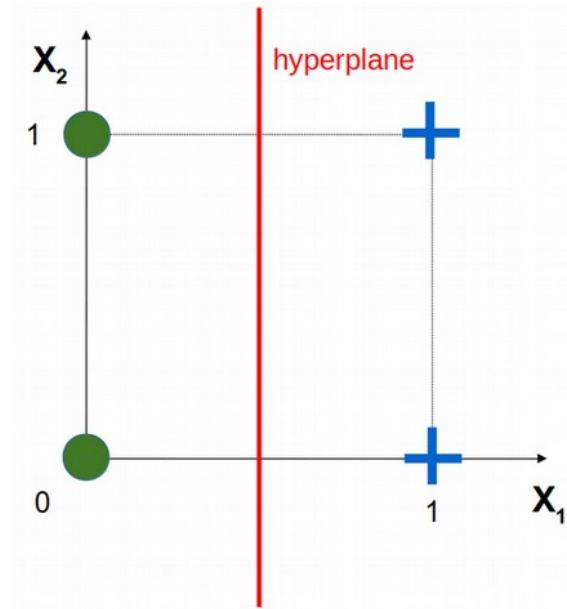
1 alpha <- vector(); beta <- vector()
2 E <- vector(); dEdalpha <- vector(); dEdbeta <- vector()
3 eta <- 0.01; alpha[1] <- 1; beta[1] <- 1 # initialize alpha and beta
4 for(i in 1:1000)
5 {
6   E[i] <- (1/n) * sum((y - alpha[i] - beta[i] * x)^2)
7   dEdalpha[i] <- -sum(2 * (y - alpha[i] - beta[i] * x)) / n
8   dEdbeta[i] <- -sum(2 * x * (y - alpha[i] - beta[i] * x)) / n
9
10  alpha[i+1] <- alpha[i] - eta * dEdalpha[i]
11  beta[i+1] <- beta[i] - eta * dEdbeta[i]
12 }
13 print(paste0("alpha = ", tail(alpha, 1), ", beta = ", tail(beta, 1)))

```

[1] "alpha = 2.89719694937354, beta = 1.94752837381973"



Let us now reconstruct the intercept and slope from gradient descent



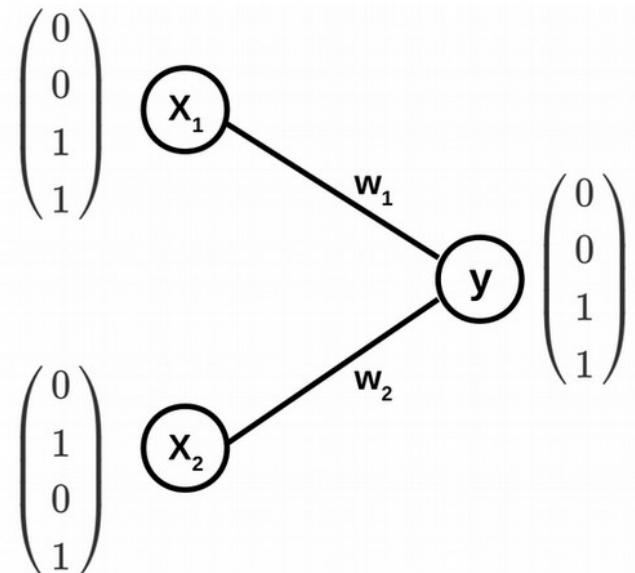
x_1	x_2	d (true) y (pred)
0	0	0 - circle
0	1	0 - circle
1	0	1 - cross
1	1	1 - cross

```

1 d <- c(0, 0, 1, 1) # true labels
2 x1 <- c(0, 0, 1, 1) # input variable x1
3 x2 <- c(0, 1, 0, 1) # input variable x2
4
5 data.frame(x1 = x1, x2 = x2, d = d)

```

x1	x2	d
0	0	0
0	1	0
1	0	1
1	1	1



$$y(w_1, w_2) = \phi(w_1 x_1 + w_2 x_2)$$

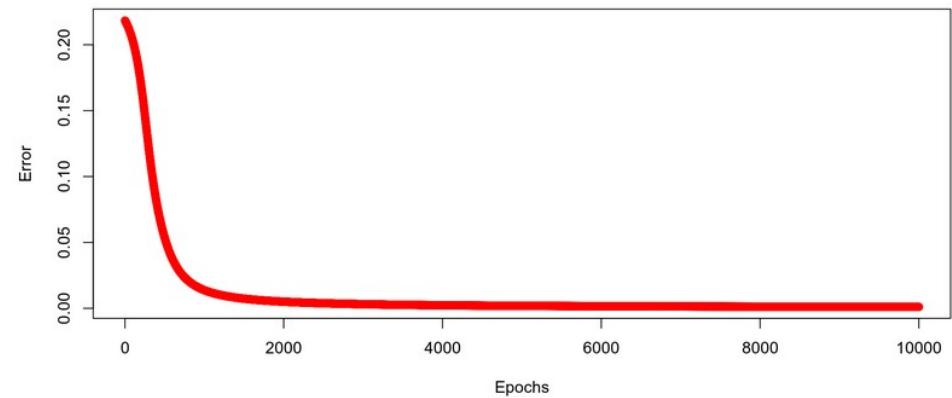
$$\phi(s) = \frac{1}{1 + e^{-s}} - \text{sigmoid}$$

$$\phi'(s) = \phi(s)(1 - \phi(s))$$

```

1 phi <- function(x){return(1/(1 + exp(-x)))} # activation function
2
3 mu <- 0.1; N_epochs <- 10000
4 w1 <- 0.1; w2 <- 0.5; E <- vector()
5 for(epochs in 1:N_epochs)
6 {
7   #Forward propagation
8   y <- phi(w1 * x1 + w2 * x2 - 3) # we use a fixed bias -3
9
10  #Backward propagation
11  E[epochs] <- (1 / (2 * length(d))) * sum((d - y)^2)
12  dE_dw1 <- - (1 / length(d)) * sum((d - y) * y * (1 - y) * x1)
13  dE_dw2 <- - (1 / length(d)) * sum((d - y) * y * (1 - y) * x2)
14  w1 <- w1 - mu * dE_dw1
15  w2 <- w2 - mu * dE_dw2
16 }
17 plot(E ~ seq(1:N_epochs), xlab="Epochs", ylab="Error", col="red")

```



$$E(w_1, w_2) = \frac{1}{2N} \sum_{i=1}^N (d_i - y_i(w_1, w_2))^2$$

$$w_{1,2} = w_{1,2} - \mu \frac{\partial E(w_1, w_2)}{\partial w_{1,2}}$$

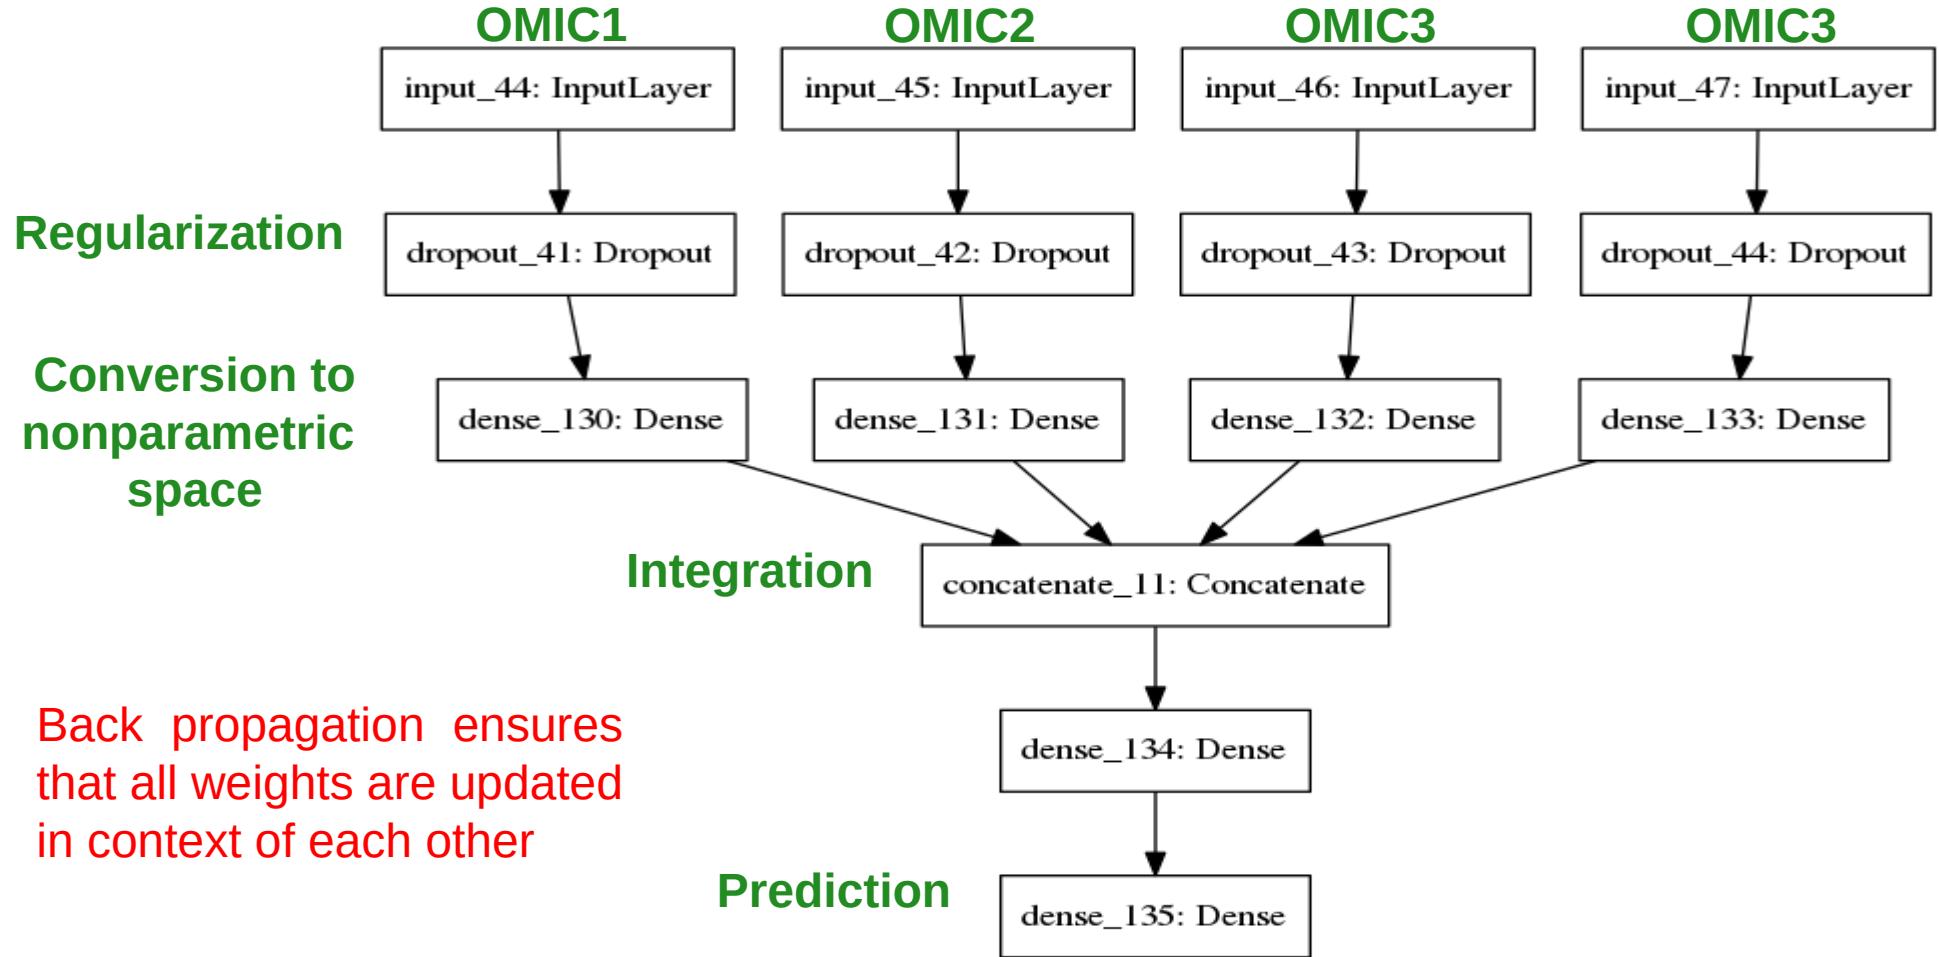
$$\frac{\partial E}{\partial w_1} = -\frac{1}{N} \sum_{i=1}^N (d_i - y_i) * y_i * (1 - y_i) * x_{1i}$$

$$\frac{\partial E}{\partial w_2} = -\frac{1}{N} \sum_{i=1}^N (d_i - y_i) * y_i * (1 - y_i) * x_{2i}$$

```
1 y
```

```
[1] 0.04742587 0.05752359 0.95730271 0.96489475
```

We nearly reconstruct true labels $d = (0, 0, 1, 1)$





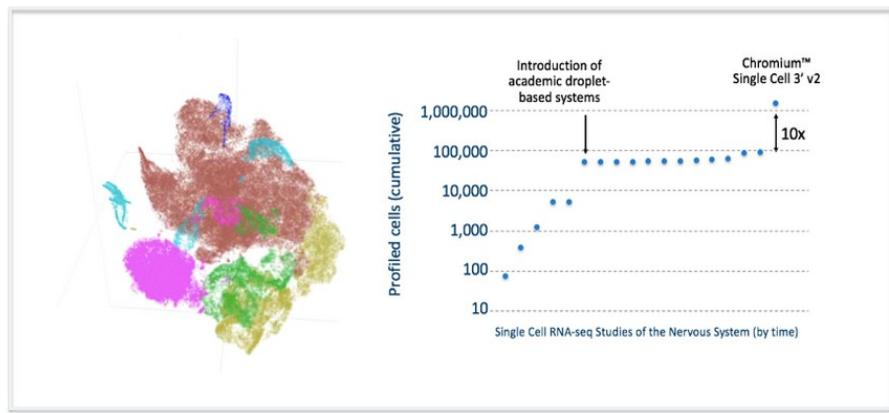
ELSEVIER

Available online at www.sciencedirect.com

ScienceDirect

[« Back to Blog](#)

< Newer Article Older Article >



Our 1.3 million single cell dataset is ready 0 KUDOS



POSTED BY grace-10x on Feb 21, 2017 at 2:28 PM

At ASHG last year, we announced our 1.3 Million Brain Cell Dataset, which is, to date, the largest dataset published in the single cell RNA-sequencing (scRNA-seq) field. Using the Chromium™ Single Cell 3' Solution (v2 Chemistry), we were able to sequence and profile 1,308,421 individual cells from embryonic mice brains. Read more in our application note [Transcriptional Profiling of 1.3 Million Brain Cells with the Chromium™ Single Cell 3' Solution](#).

Single cells make big data: New challenges and opportunities in transcriptomics

Philipp Angerer¹, Lukas Simon¹, Sophie Tritschler¹, F. Alexander Wolf¹, David Fischer¹ and Fabian J. Theis^{1,2}

Abstract

Recent technological advances have enabled unprecedented insight into transcriptomics at the level of single cells. Single cell transcriptomics enables the measurement of transcriptomic information of thousands of single cells in a single experiment. The volume and complexity of resulting data make it a paradigm of big data. Consequently, the field is presented with new scientific and, in particular, analytical challenges where currently no scalable solutions exist. At the same time, exciting opportunities arise from increased resolution of single-cell RNA sequencing data and improved statistical power of ever growing datasets. Big single cell RNA sequencing data promises valuable insights into cellular heterogeneity which may significantly improve our understanding of biology and human disease. This review focuses on single cell transcriptomics and highlights the inherent opportunities and challenges in the context of big data analytics.

Addresses

¹ Institute of Computational Biology, Helmholtz Center Munich, German Research Center for Environmental Health, Neuherberg, Germany

² Department of Mathematics, Technical University of Munich, Garching, Germany

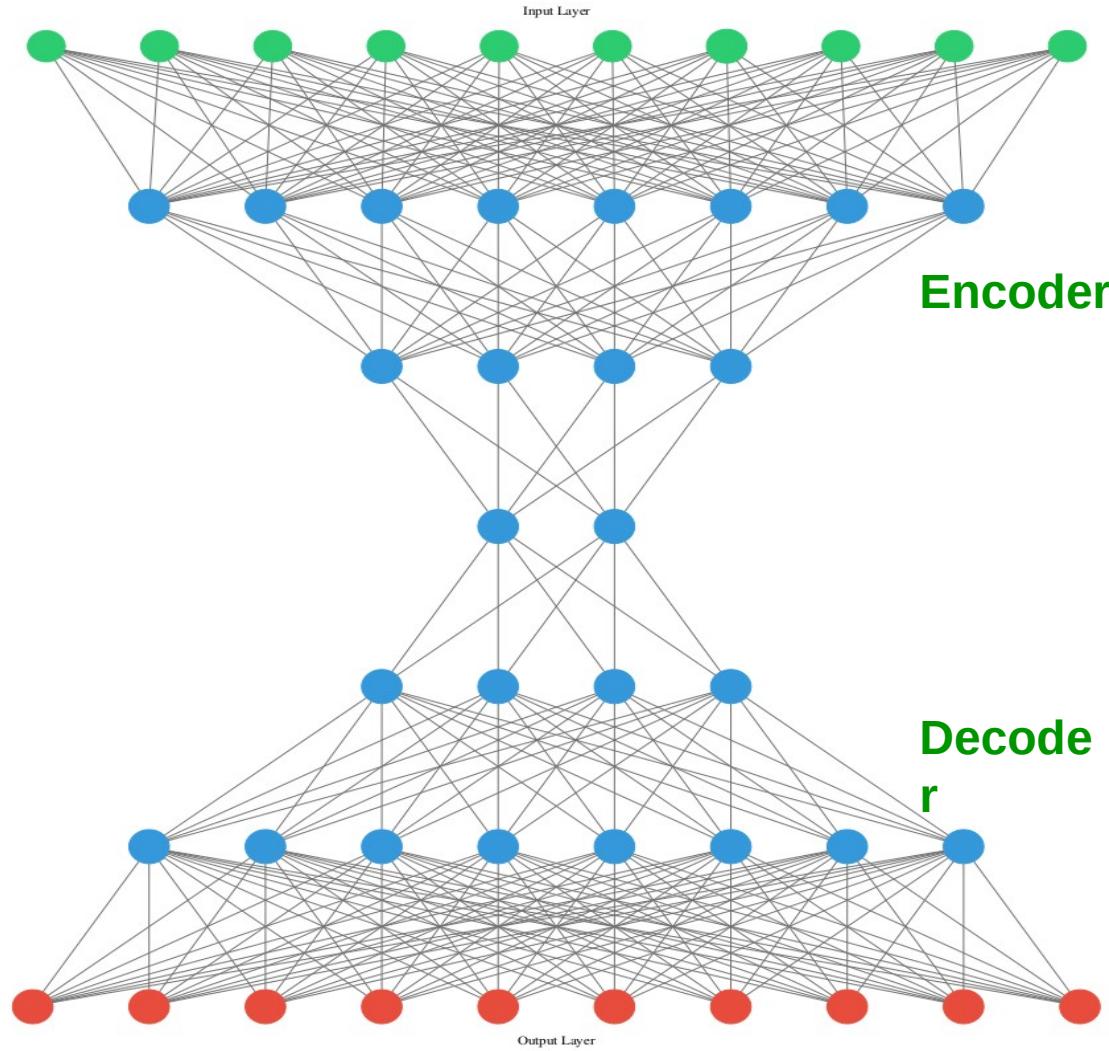
Corresponding author: Theis, Fabian J. Institute of Computational Biology, Helmholtz Center Munich, German Research Center for Environmental Health, Neuherberg, Germany. (fabian.theis@helmholtz-muenchen.de)

Current Opinion in Systems Biology 2017, 4:85–91

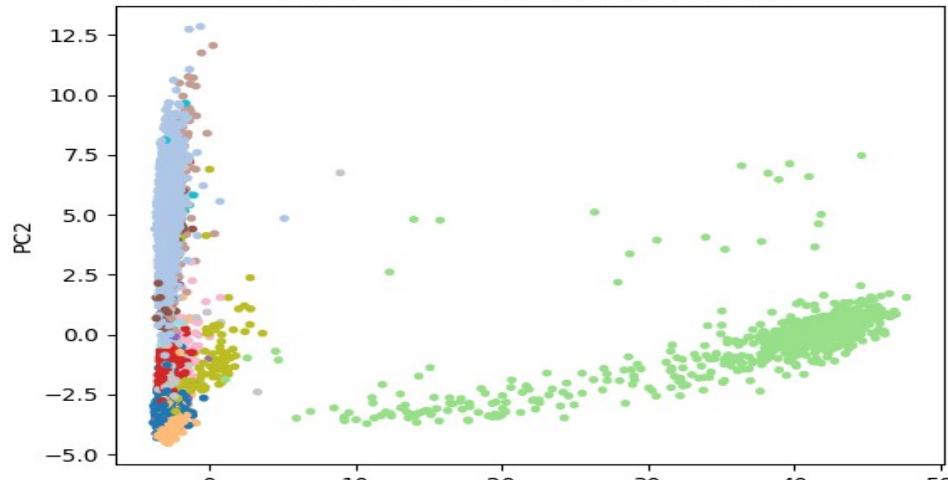
This review comes from a themed issue on **Big data acquisition and**

volume – the amount of data, velocity – the required processing speed, veracity – trustworthiness and availability, and variety – necessary model complexity [2]. The traditional scientific big data field is astronomy because of the huge *volume* of image data produced by telescopes with a high daily *velocity* [3]. Big data has also reached biology, mainly driven through the advent of next generation sequencing technology. For biologists, assessing *veracity* through statistical means is nothing new.

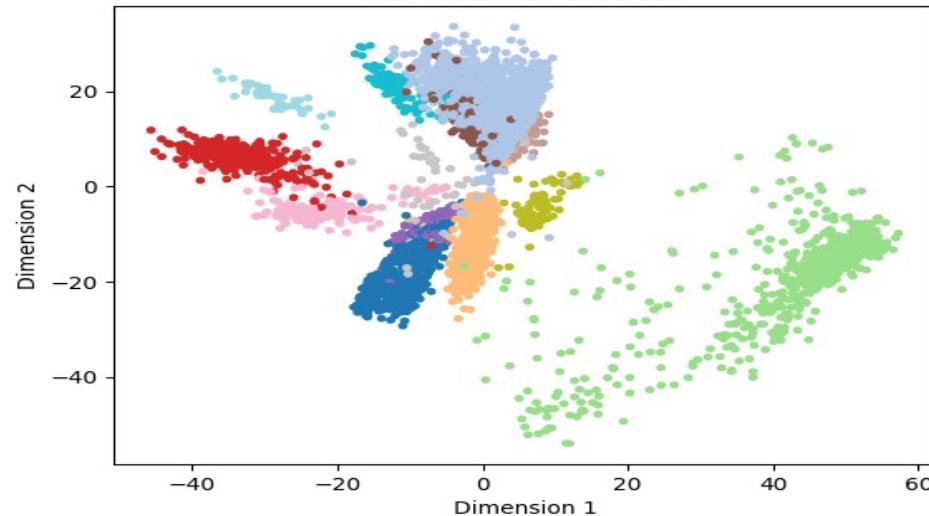
Recent technological advances now allow the profiling of single cells at a *variety* of omic layers (genomes, epigenomes, transcriptomes and proteomes) at an unprecedented level of resolution [4]. Single cell transcriptomics (SCT) entails the profiling of all messenger RNAs present in a single cell and constitutes the most widely-used sc profiling technology [4]. Unlike bulk RNA-seq profiling where sequencing libraries are generated from thousands of cells, scRNA-seq technologies isolate single cells and generate cell-specific sequencing libraries (e.g. Fluidigm [5]) mark RNA content with a cell-specific molecular barcode [6–9]. Both approaches generate gene expression estimates at the single cell level [10]. SCT enables, for the first time, the measurement of the transcriptomic information of thousands, and up to millions of single cells, in a single experiment [7]. The complexity of SCT data coupled with the massive volume inherent to next generation sequencing data makes it a paradigm of big data.



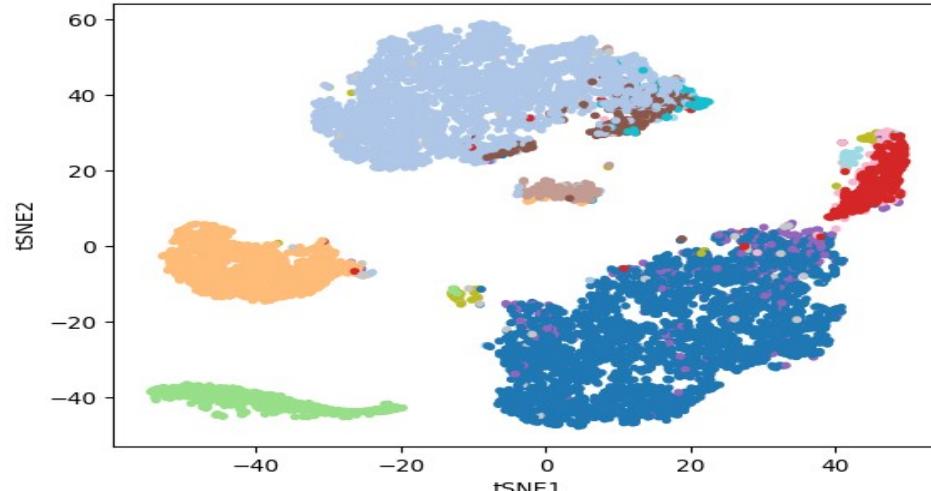
Principal Component Analysis (PCA)



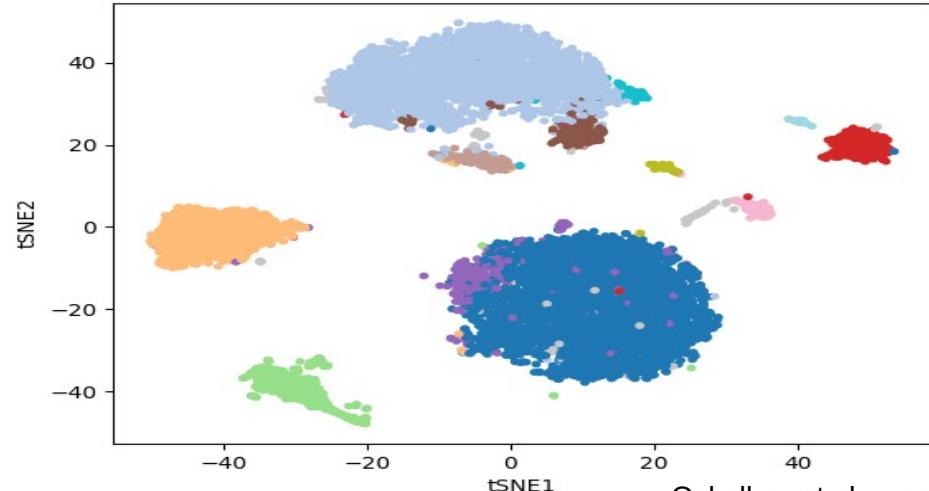
Autoencoder: 8 Layers



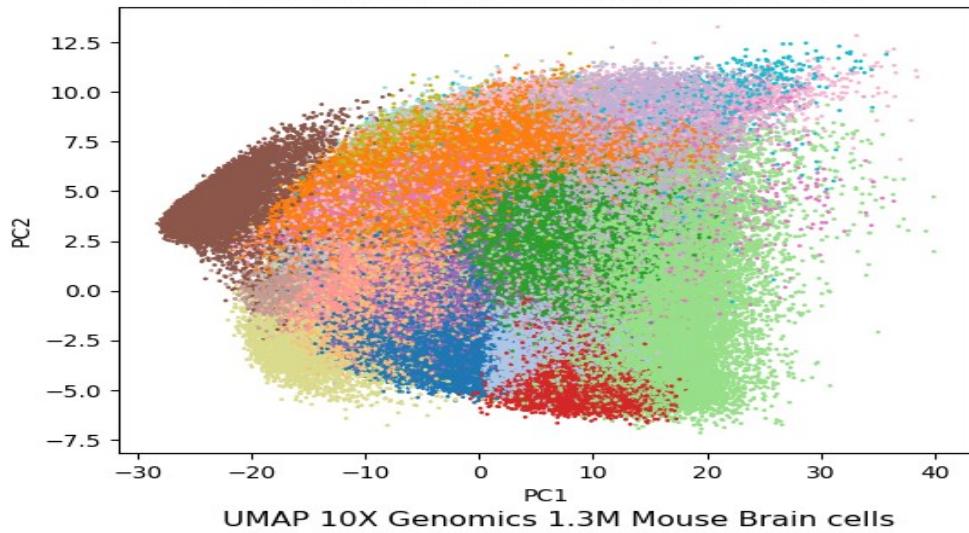
tSNE on PCA



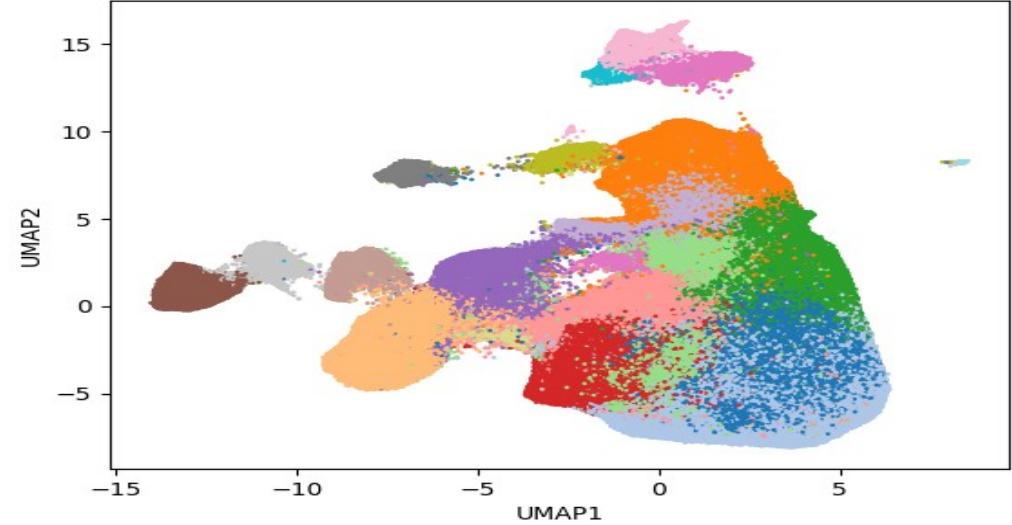
tSNE on Autoencoder: 8 Layers



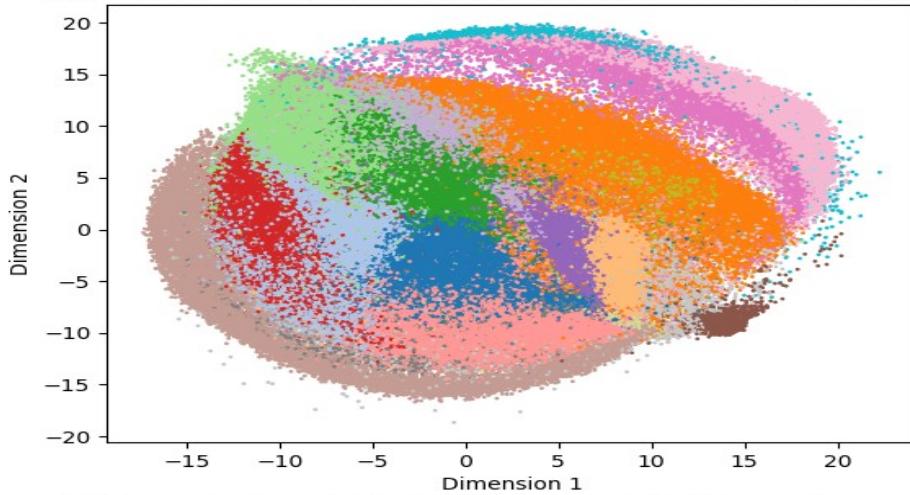
PCA, 10X Genomics 1.3M Mouse Brain Cells



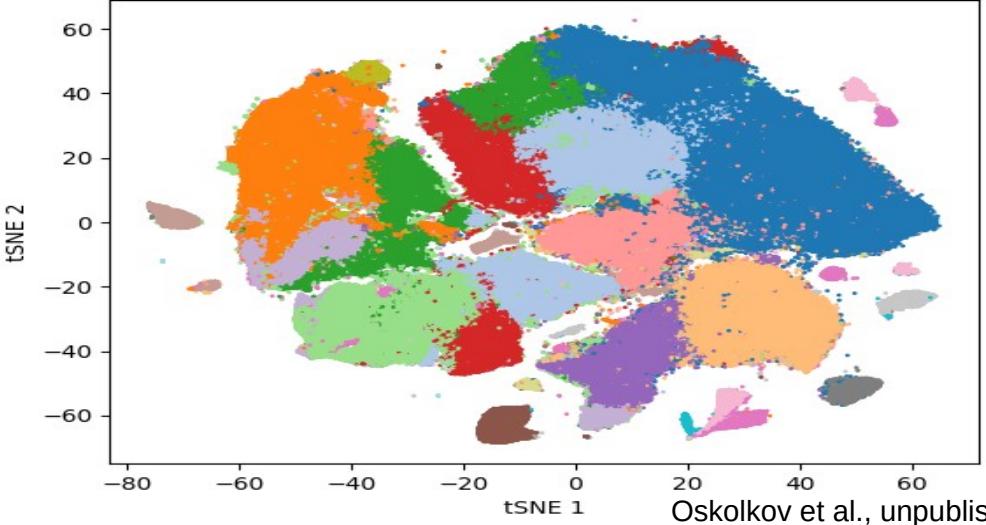
UMAP 10X Genomics 1.3M Mouse Brain cells



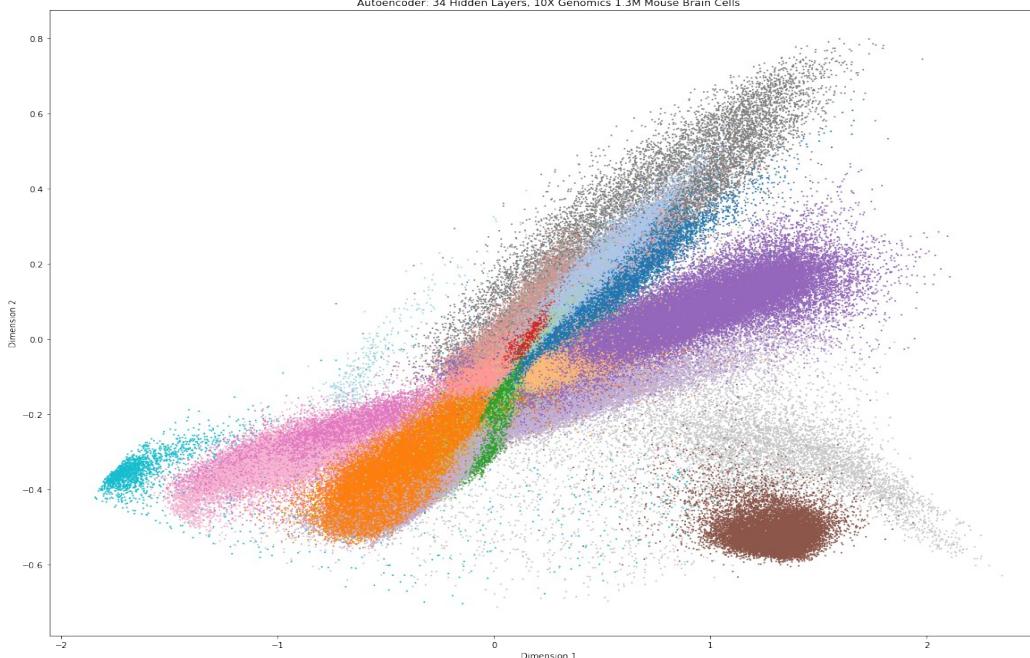
Autoencoder 10 Hidden Layers, 10X Genomics 1.3M Mouse Brain cells



tSNE perplexity = 350, 10X Genomics 1.3M Mouse Brain cells

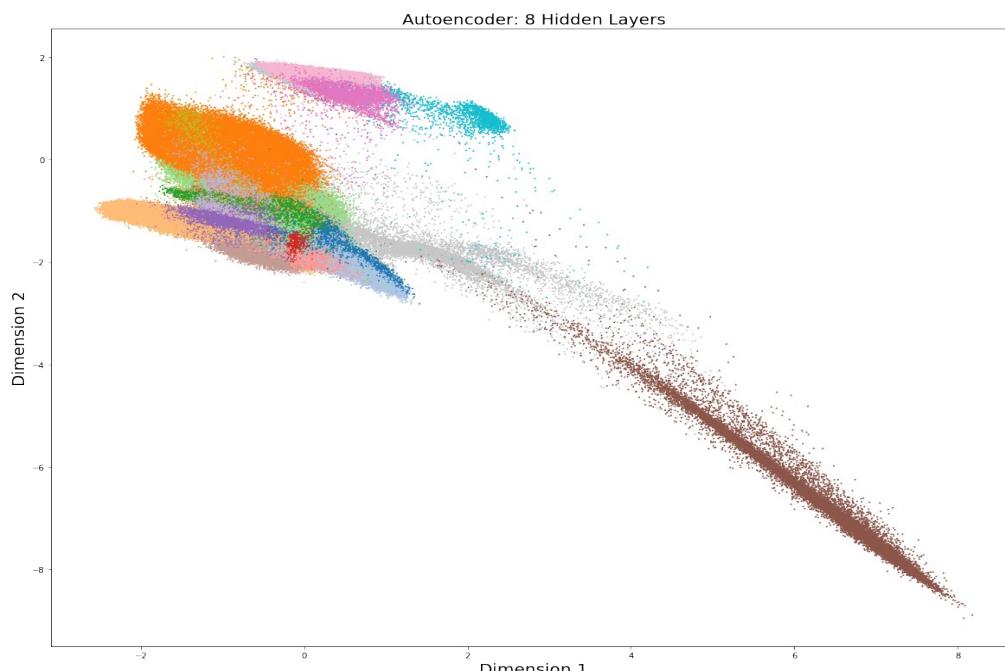


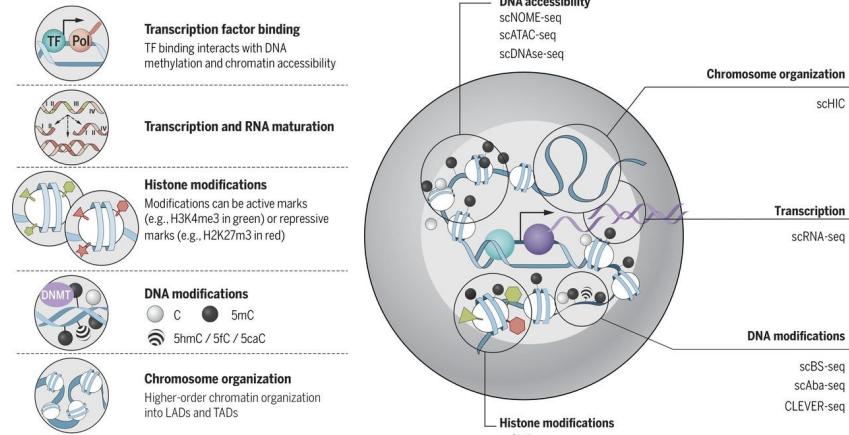
Oskolkov et al., unpublished



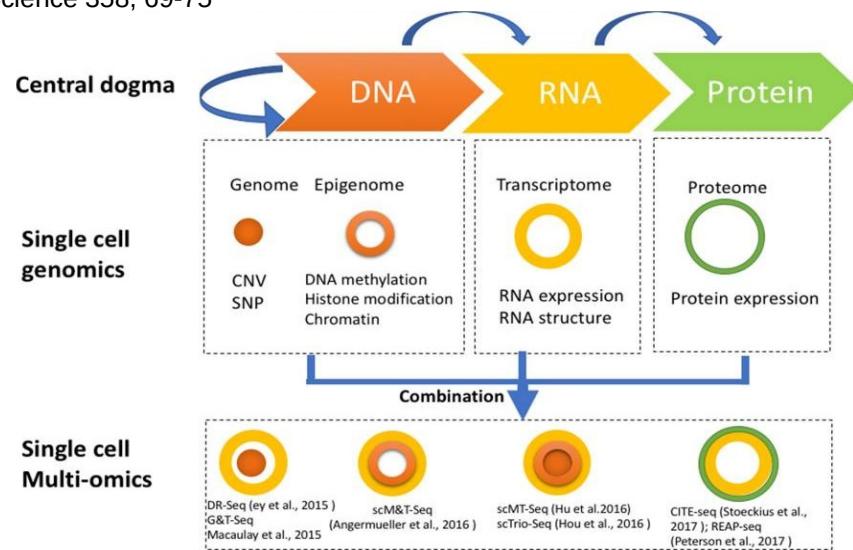
Autoencoders themselves are perhaps not optimal for visualization of scOmics

Autoencoders can be promising for non-linear data pre-processing, the bottleneck can potentially be fed to tSNE / UMAP

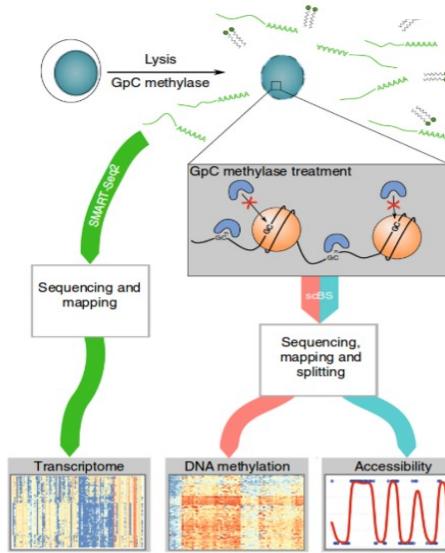




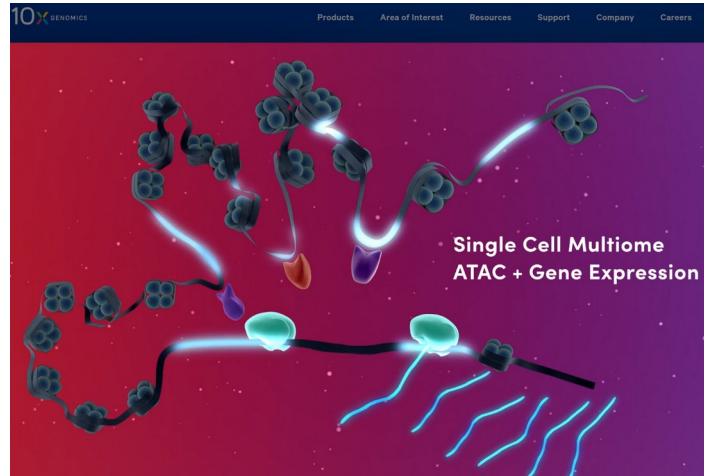
Kelsey et al., 2017, Science 358, 69-75



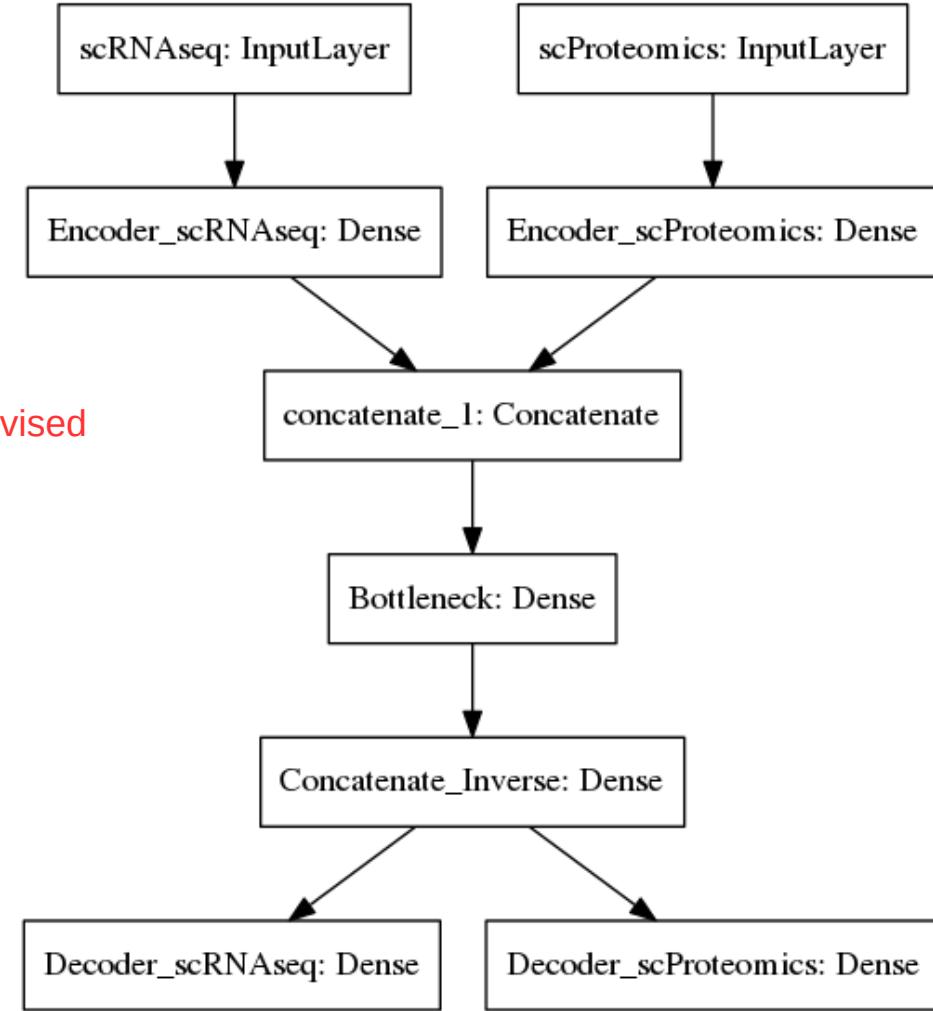
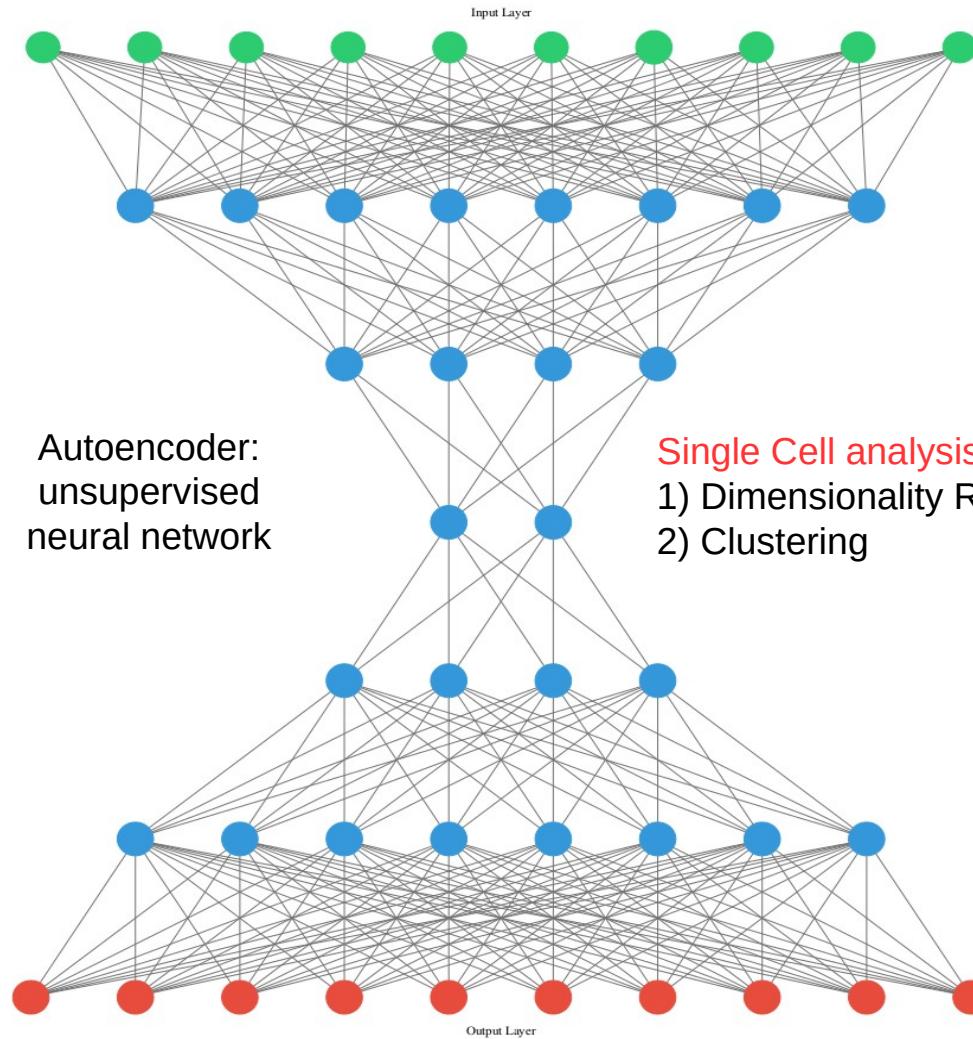
Hu et al., 2018, Frontier in Cell and Developmental Biology 6, 1-13

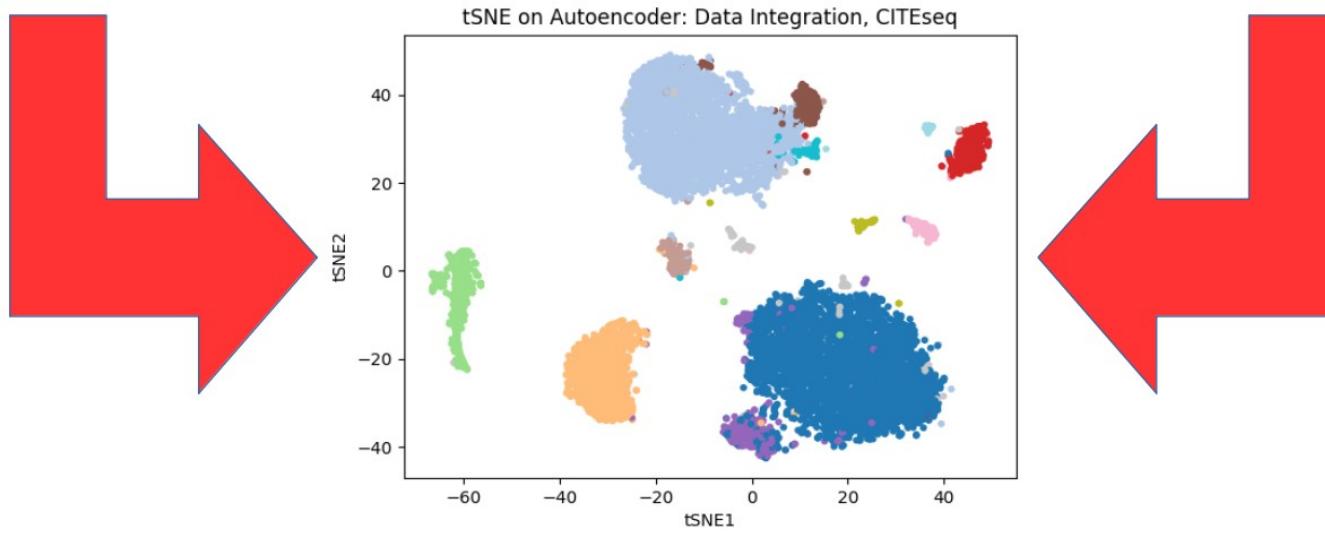
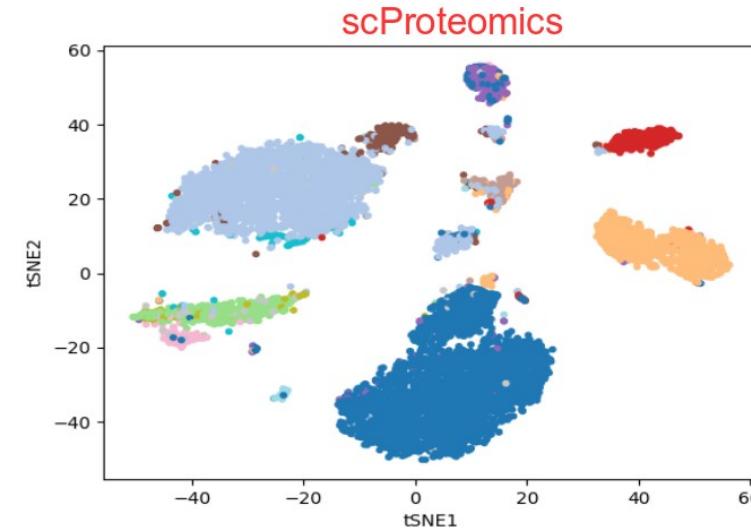
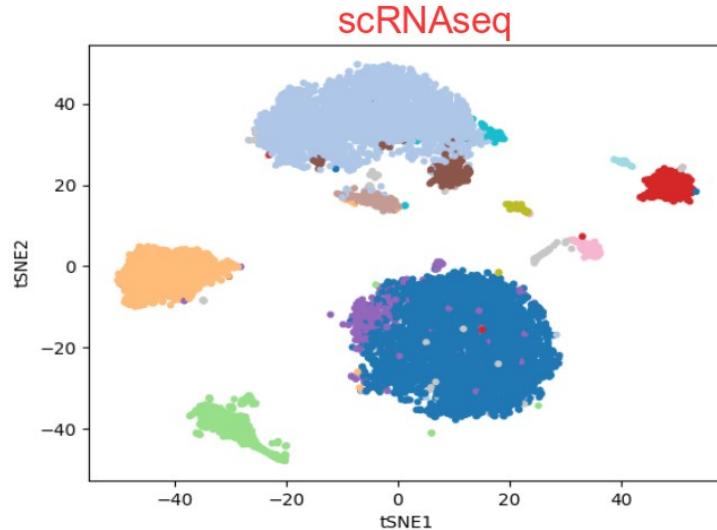


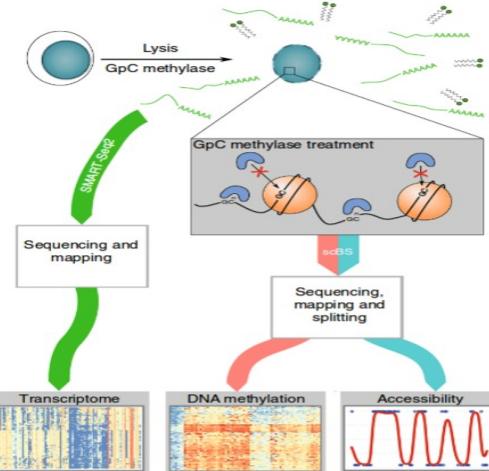
Clark et al., 2018, Nature Communications 9, 781



Autoencoder for data integration

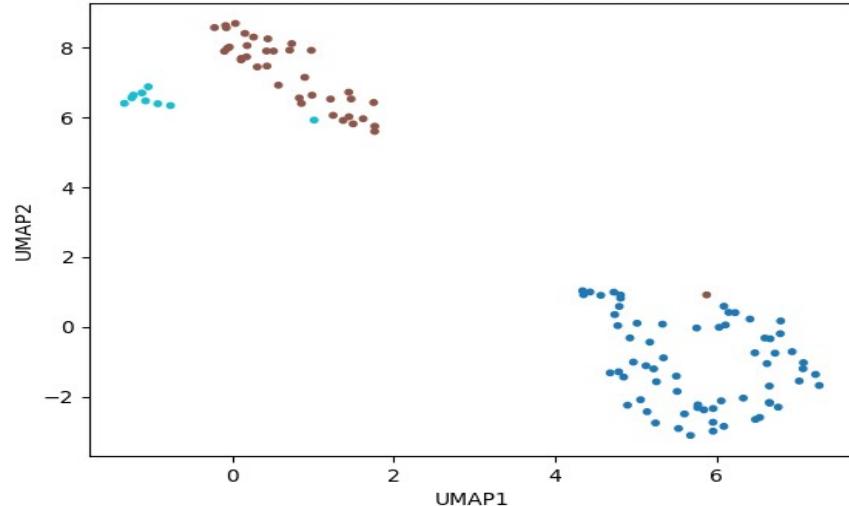




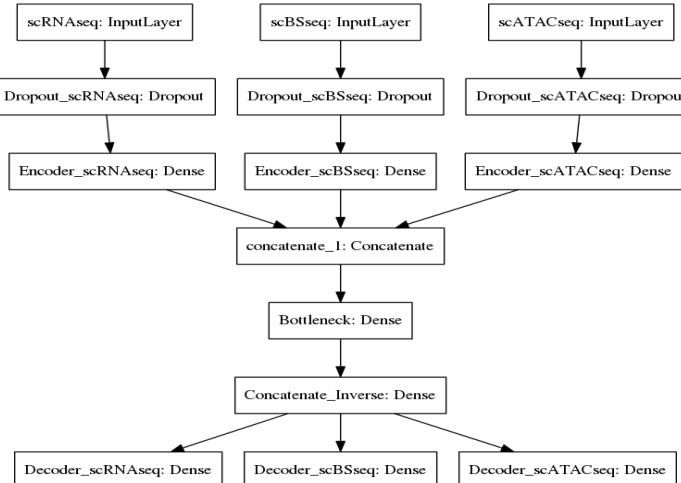


scNMTseq: Clark et al., 2018, Nature Communications 9, 781

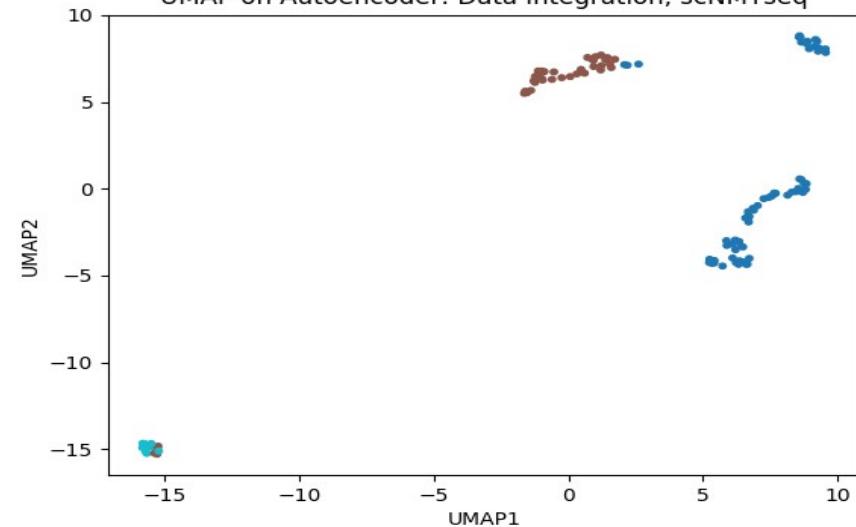
UMAP on PCA: scNMTseq, scRNAseq

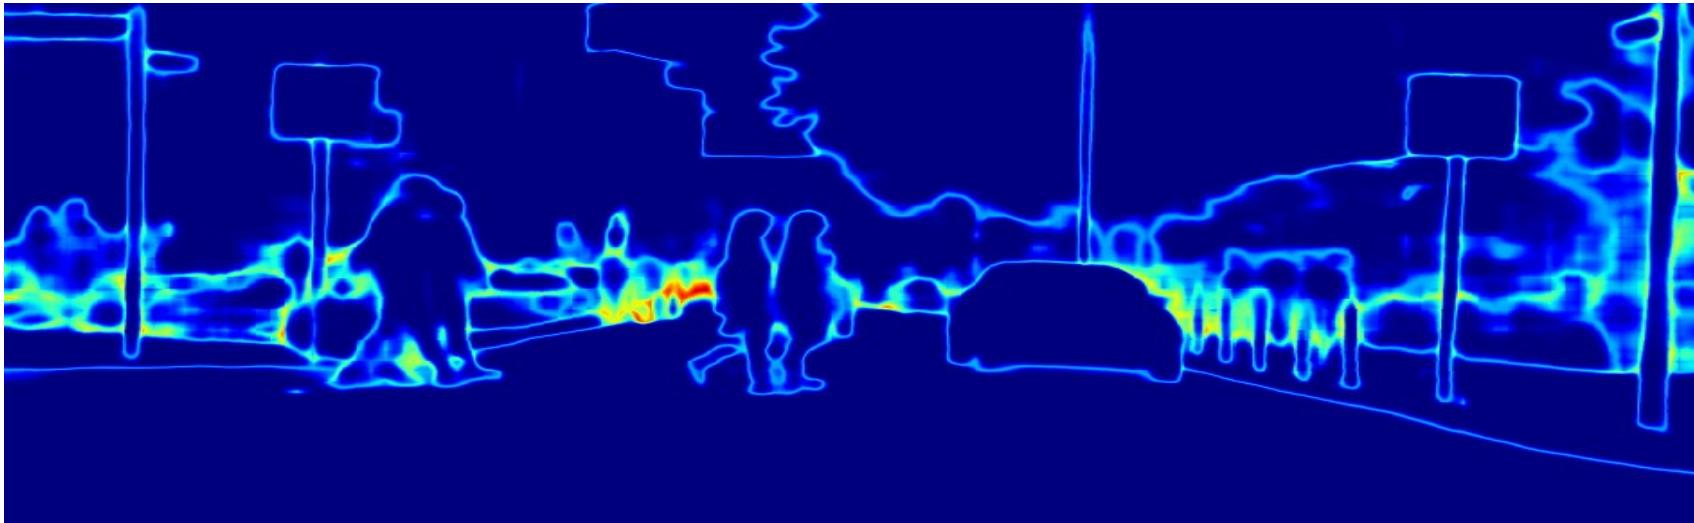


From Single
To
multi-Omics

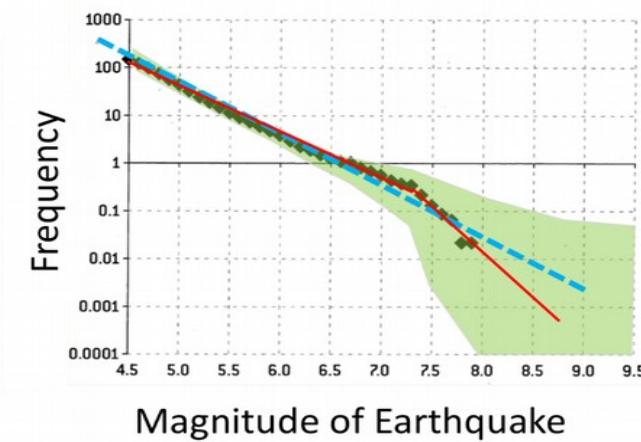
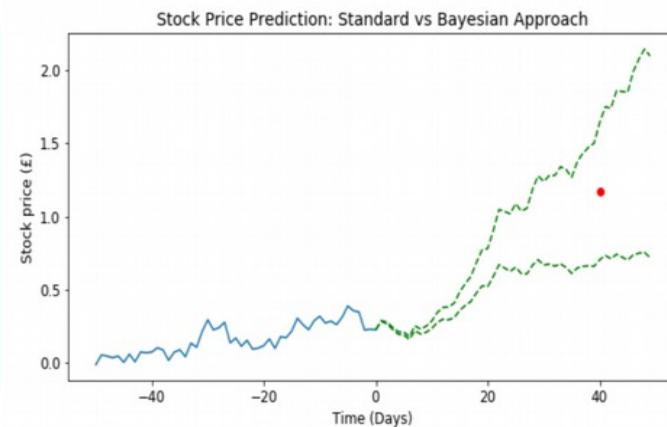


UMAP on Autoencoder: Data Integration, scNMTseq

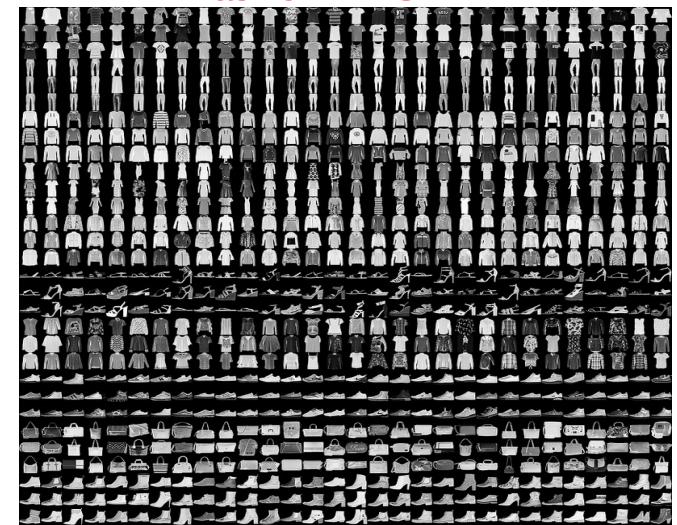




Intelligence is to know how much you do not know



Frequentist image recognition



```
In [24]: # normalize inputs from 0-255 to 0-1.0
X_train = X_train.reshape(X_train.shape[0], 1, 28, 28).astype('float32')
X_test = X_test.reshape(X_test.shape[0], 1, 28, 28).astype('float32')
X_train = X_train / 255.0
X_test = X_test / 255.0

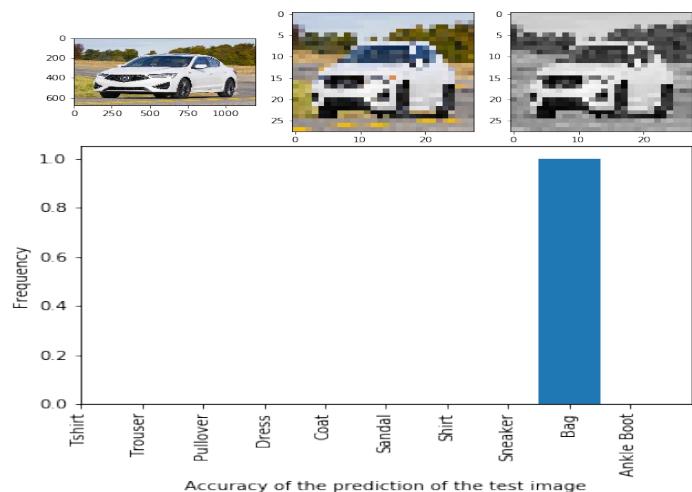
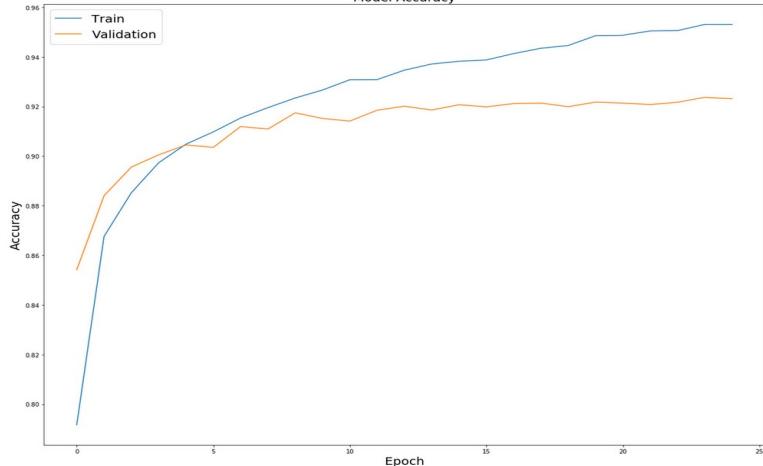
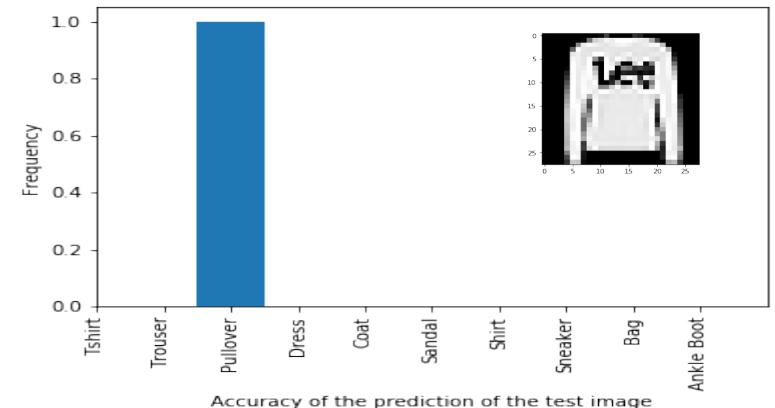
In [25]: # one hot encode outputs
y_train = np_utils.to_categorical(y_train)
y_test = np_utils.to_categorical(y_test)
num_classes = y_train.shape[1]
print(num_classes)
10

In [27]: # Create the model
model = Sequential()
model.add(Conv2D(32, (3, 3), input_shape=(1, 28, 28), padding='same', activation='relu',
                kernel_constraint=maxnorm(3)))
model.add(Dropout(0.2))
model.add(Conv2D(32, (3, 3), padding='same', activation='relu',
                kernel_constraint=maxnorm(3)))
model.add(MaxPooling2D(pool_size=(2, 2)))
model.add(Flatten())
model.add(Dense(128, activation='relu', kernel_constraint=maxnorm(3)))
model.add(Dense(10, activation='softmax'))

# Compile the model
model.compile(optimizer='sgd', loss='categorical_crossentropy', metrics=['accuracy'])
print(model.summary())

Layer (type)                 Output Shape              Param #
=====================================================================
conv2d_8 (Conv2D)            (None, 32, 28, 28)      320
dropout_7 (Dropout)          (None, 32, 28, 28)      0
conv2d_9 (Conv2D)            (None, 32, 28, 28)      9248
max_pooling2d_4 (MaxPooling2D) (None, 32, 14, 14)     0
flatten_4 (Flatten)          (None, 6272)             0
dense_7 (Dense)              (None, 512)              3211776
dropout_8 (Dropout)          (None, 512)              0
dense_8 (Dense)              (None, 10)               5130
=====
Total params: 3,226,474
Trainable params: 3,226,474
Non-trainable params: 0
None

In [28]: # Fit the model
history = model.fit(X_train, y_train, validation_data=(X_test, y_test), epochs=epochs, batch_size=32,
                     history = model.fit(X_train, y_train, epochs = epochs, verbose = 1, validation_split = 0.25,
```



PyMC3, Edward, TensorFlow Probability

```
In [8]: x_train = x_train.reshape((x_train.shape[0],D))
x_test = x_test.reshape((x_test.shape[0],D))
print(x_train.shape)
print(x_test.shape)
(50000, 784)
(10000, 784)

In [9]: from keras.utils import to_categorical
y_train = to_categorical(y_train)
y_test = to_categorical(y_test)
print(y_train.shape)
print(y_test.shape)
(50000, 10)
(10000, 10)

In [10]: ed.set_seed(314159)
# D = 784 # number of images in a minibatch.
D = 784 # number of features.
K = 10 # number of classes.

# Create a placeholder to hold the data (in minibatches) in a TensorFlow graph.
x = tf.placeholder(tf.float32, [None, D])
# Normal distribution for the variational. Note that the syntax assumes TensorFlow 1.2.
w = Normal(loc=tf.zeros([D, K]), scale=tf.ones([D, K]))
# Normal(loc=tf.zeros([K]), scale=tf.ones([K]))
# Categorical likelihood for classification.
y = Categorical(tf.matmul(w, x))

In [11]: # Construct the q(w) and q(b). In this case we assume Normal distributions.
q_w = Normal(loc=tf.Variable(tf.random_normal([D, K])),
              scale=tf.nn.softplus(tf.Variable(tf.random_normal([D, K]))))
q_b = Normal(loc=tf.Variable(tf.random_normal([K])),
              scale=tf.nn.softplus(tf.Variable(tf.random_normal([K]))))

In [12]: def generator(arrays, batch_size=N):
    parts = [0] * len(arrays) # pointers to where we are in iteration
    batches = []
    for i, array in enumerate(arrays):
        start = starts[i]
        stop = start + batch_size
        diff = array.shape[0]
        if diff <= 0:
            start = np.array([start])
            stop = np.array([stop])
            starts[i] = batch_size
        else:
            batch = np.concatenate((array[start:], array[:diff]))
            starts[i] = diff
        batches.append(batch)
    yield batches
    cifar10 = generator([x_train, y_train], N)

In [13]: # We use a placeholder for the labels in anticipation of the training data.
y_ph = tf.placeholder(tf.int32, [N])
# Define the VI inference technique, ie. minimize the KL divergence between q and p.
inf_type = 'KLqp'
inf_args = {}
inf_options = {}
infiagnostics = {}
inf_inference = inference.inference(n_iter=50000, n_print=100, scale=(float(x_train.shape[0]) / N))
inf_options['n_iter'] = 50000
inf_options['n_print'] = 100
inf_options['scale'] = float(x_train.shape[0]) / N
sess = tf.InteractiveSession()
# Initialize all the variables in the session.
tf.global_variables_initializer().run()
# Let the training begin. We load the data in minibatches and update the VI inference using each new batch.
for i in range(inference.n_iter):
    X_batch = cifar10.next_batch(100)
    y_batch = np.argmax(Y_batch, axis=1)
    infiagnostics = inf_inference.update(feed_dict={x: X_batch, y_ph: y_batch})
    inf_options['progressive'] = infiagnostics['progressive']
    inference.print_progressive(**inf_options)
    sess.run(tf.global_variables_initializer())
    # Compute the accuracy of the model.
    # For each sample we compute the predicted class and compare with the test labels.
    # Predicted class is defined as the one which has maximum probability.
    # Finally we make a histogram of accuracies for the test data.
    accy_test = []
    for prob in prob_list:
        y_trn_prct = np.argmax(prob, axis=1).astype(np.float32)
        y_tst_prct = np.argmax(y_test, axis=1).mean() * 100
        accy_test.append(accy)

    plt.hist(accy_test)
    plt.title("Histogram of prediction accuracies in the CIFAR10 test data")
    plt.xlabel("Accuracy")
    plt.ylabel("Frequency")
    plt.show()

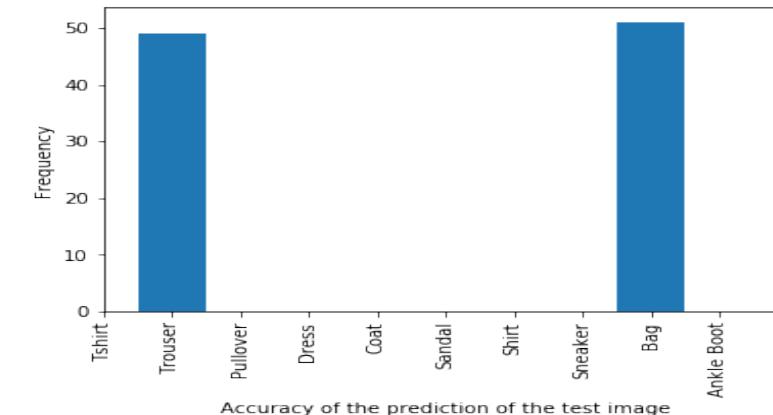
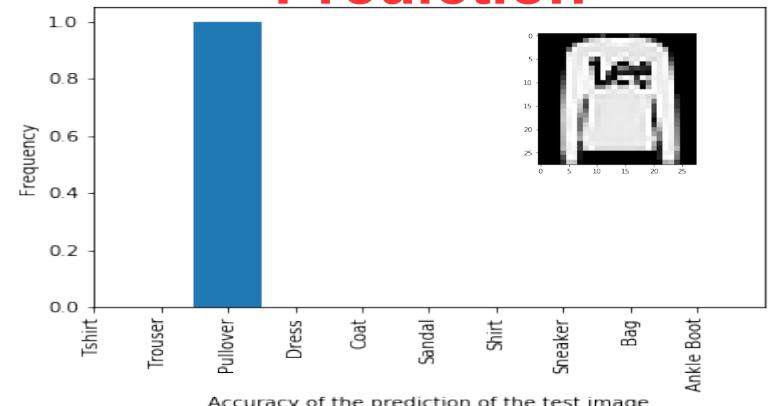
Elapsed: 221s | Loss: 85453.266
```

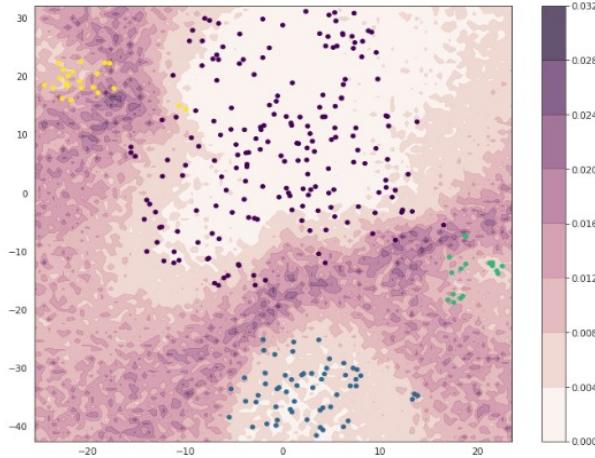
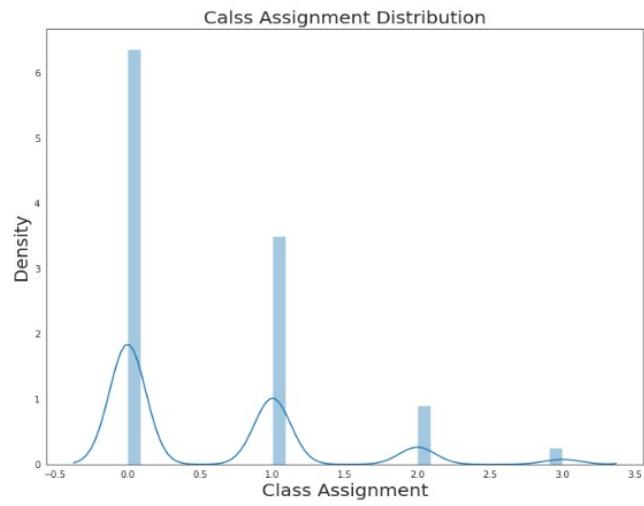
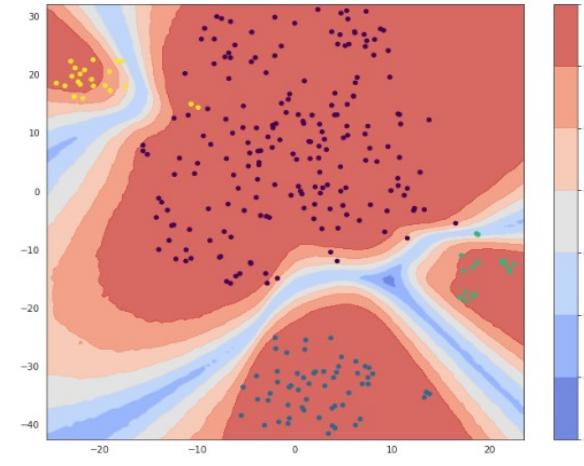
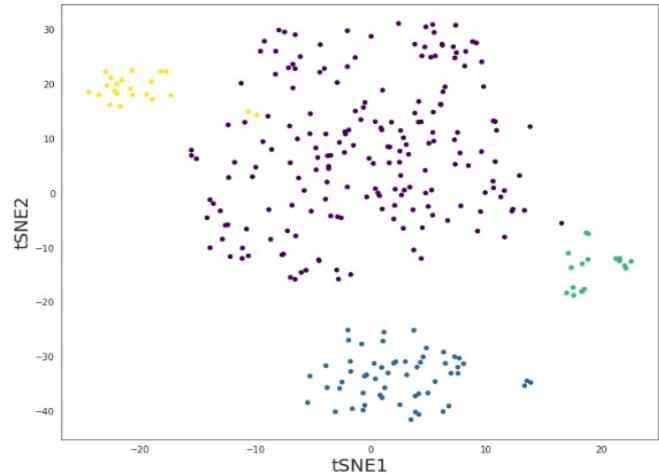
```
In [14]: n_samples = 100
prob_list = []
samples = []
w_samples = []
b_samples = []
for i in range(n_samples):
    w_samp = q_w.sample()
    b_samp = q_b.sample()
    samples.append(w_samp)
    w_samples.append(w_samp)
    b_samples.append(b_samp)
    # also compute the probability of each class for each (w,b) sample.
    prob = tf.nn.softmax(tf.matmul(x_test, w_samp) + b_samp)
    prob_stacked = tf.stack(prob, axis=1)
    sample = tf.concat([tf.reshape(w_samp, [-1]), b_samp], 0)
    samples.append(sample.eval())

In [15]: # Compute the accuracy of the model.
# For each sample we compute the predicted class and compare with the test labels.
# Predicted class is defined as the one which has maximum probability.
# Finally we make a histogram of accuracies for the test data.
accy_test = []
for prob in prob_list:
    y_trn_prct = np.argmax(prob, axis=1).astype(np.float32)
    y_tst_prct = np.argmax(y_test, axis=1).mean() * 100
    accy_test.append(accy)

plt.hist(accy_test)
plt.title("Histogram of prediction accuracies in the CIFAR10 test data")
plt.xlabel("Accuracy")
plt.ylabel("Frequency")
plt.show()
```

Prediction





Bartoschek, Oskolkov et al., Nature Communications 2018

Deep generative modeling for single-cell transcriptomics

Romain Lopez¹, Jeffrey Regier¹, Michael B. Cole², Michael I. Jordan^{1,3} and Nir Yosef^{1,4,5*}

Single-cell transcriptome measurements can reveal unexplored biological diversity, but they suffer from technical noise and bias that must be modeled to account for the resulting uncertainty in downstream analyses. Here we introduce single-cell variational inference (scVI), a ready-to-use scalable framework for the probabilistic representation and analysis of gene expression in single cells (<https://github.com/YosefLab/scVI>). scVI uses stochastic optimization and deep neural networks to aggregate information across similar cells and genes and to approximate the distributions that underlie observed expression values, while accounting for batch effects and limited sensitivity. We used scVI for a range of fundamental analysis tasks including batch correction, visualization, clustering, and differential expression, and achieved high accuracy for each task.

nature methods

Article

<https://doi.org/10.1038/s41592-024-02201-0>

scGPT: toward building a foundation model for single-cell multi-omics using generative AI

Received: 12 July 2023

Accepted: 30 January 2024

Published online: 26 February 2024

Check for updates

Haotian Cui^{1,2,3,8}, Chloe Wang^{1,2,3,8}, Hassaan Maan^{1,3,4}, Kuan Pang^{1,2,3}, Fengning Luo^{2,3}, Nan Duan^{1,5} & Bo Wang^{1,2,3,4,6,7}

Generative pretrained models have achieved remarkable success in various domains such as language and computer vision. Specifically, the



ARTICLE

DOI: [10.1038/s41467-018-04368-5](https://doi.org/10.1038/s41467-018-04368-5)

OPEN

Interpretable dimensionality reduction of single cell transcriptome data with deep generative models

Jiarui Ding^{1,2,3,4}, Anne Condon¹ & Sohrab P. Shah^{1,2,3,5}

Huang et al. *Genome Biology* (2023) 24:259
<https://doi.org/10.1186/s13059-023-03100-x>

Genome Biology

RESEARCH

Open Access

Evaluation of deep learning-based feature selection for single-cell RNA sequencing data analysis



Hao Huang^{1,2,3}, Chunlei Liu^{1,3}, Manoj M. Wagle^{1,2,3} and Pengyi Yang^{1,2,3,4*}

Biomedicine & Pharmacotherapy 165 (2023) 115077



Contents lists available at ScienceDirect

Biomedicine & Pharmacotherapy

journal homepage: www.elsevier.com/locate/biopha

Review

Deep learning applications in single-cell genomics and transcriptomics data analysis

Nafiseh Erfanian^a, A. Ali Heydari^{b,c}, Adib Miraki Feriz^a, Pablo Iañez^d, Afshin Derakhshani^e, Mohammad Ghasemigol^f, Mohsen Farahpour^g, Seyyed Mohammad Razavi^g, Saeed Nasseri^h, Hossein Safarpour^{h,i}, Amirhossein Sahebkar^{i,j,k,**}



TAGGED IN

DI For Life Sciences

Towards Data Science

Your home for data science. A Medium publication sharing concepts, ideas and codes.

More information

FOLLOWERS

688K

ELSEWHERE



MORE, ON MEDIUM

DI For Life Sciences



Nikolay Oskolkov in Towards Data Science

Aug 18, 2021 · 13 min read ★



DEEP LEARNING FOR LIFE SCIENCES

Deep Learning on Human Microbiome

Infer microbial...

Read more...



4 responses

Nikolay Oskolkov in Towards Data Science
May 20, 2019 · 7 min read ★



Deep Learning for Life Sciences

Deep Learning for Clinical Diagnostics

Read more...

141

Nikolay Oskolkov in Towards Data Science
Aug 6, 2019 · 6 min read ★



Deep Learning for Life Sciences

Why Biology is Sceptic Towards AI

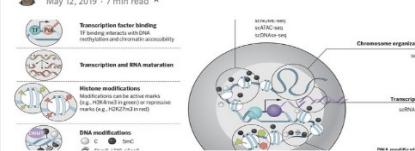
And why Precision...

Read more...

159

1 response

Nikolay Oskolkov in Towards Data Science
May 12, 2019 · 7 min read ★



Deep Learning for Life Sciences

Deep Learning for Data Integration

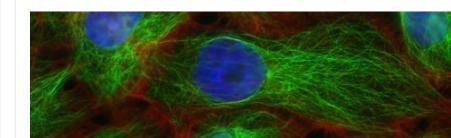
Synergistic effects...

Read more...

126

5 responses

Nikolay Oskolkov in Towards Data Science
Jul 7, 2019 · 9 min read ★



Deep Learning for Life Sciences

Deep Learning on Microscopy Imaging

Read more...

207

1 response

Nikolay Oskolkov in Towards Data Science
May 5, 2019 · 9 min read ★



Deep Learning for Life Sciences

Deep Learning for Single Cell Biology

Read more...

406

7 responses

Nikolay Oskolkov in Towards Data Science
Apr 28, 2019 · 6 min read ★



Deep Learning for Life Sciences

Deep Learning on Ancient DNA

Reconstructing the Human...

Read more...

381

4 responses



*Knut och Alice
Wallenbergs
Stiftelse*



Vetenskapsrådet



**LUNDS
UNIVERSITET**